

# Smart Agentic Task Router with Neo4j-Enhanced Agent Profiling (Updated Proposal)

Incorporating Instructor Feedback & Filling All Gaps

**Student:** Azad Jagtap **Domain:** Multi-Agent Systems & Knowledge Representation

---

## 1. Problem Statement

Modern AI workflows increasingly depend on **specialized agents**—such as web searchers, code analyzers, summarizers, and data visualizers—each optimized for distinct tasks. Users, however, often struggle to determine which agent best fits their query. This leads to:

- Suboptimal results
- Inefficient resource usage
- Poor user experience

Existing routing approaches (keyword matching, intent classification) fail to model the **nuanced relationships** between:

- Task requirements
- Agent capabilities
- Performance histories
- Complexity levels
- Agent interoperability

**Core Problem:**

*How do we intelligently route user queries to the most appropriate specialized agent using a knowledge graph that supports explainability, adaptability, and continuous learning?*

---

## 2. Proposed Solution: Smart Agentic Router

The Smart Agentic Router uses **Neo4j**, **RDF/OWL modeling**, **SHACL constraints**, and a **multi-agent architecture** to:

1. Model agent capabilities, relationships, fallback logic
2. Understand task types and query semantics
3. Execute SPARQL/Cypher queries to rank and select the optimal agent
4. Learn over time using performance feedback and routing outcomes

The system provides **explainable routing** by exposing the graph traversal path used in each decision.

---

## 3. Knowledge Graph Design (Expanded with Instructor Feedback)

### 3.1 Core Classes (RDF/OWL)

**Agent Classes**

- **ex:Agent**: Generic agent

- **ex:SpecializedAgent** ← ex:Agent
- **ex:RouterAgent** ← ex:Agent

## Task & Query Classes

- ex:Task
- ex:TaskType
- ex:Query

## Capability & Performance Classes

- ex:Capability
  - ex:RoutingDecision
  - ex:PerformanceRecord
- 

## 3.2 Key Relationships

### Agent-Capability

- ex:hasCapability
- ex:capabilityLevel (0.0-1.0)
- ex:fallbackAgent
- ex:similarTo

### Task-Capability

- ex:requiresCapability
- ex:complexityLevel

### Query-Agent

- ex:routedTo
- ex:confidence

### Performance

- ex:successfullyHandled
  - ex:failureCount
  - ex:latency
- 

## 3.3 Additional RDF Properties (Added per Feedback)

- ex:inputFormat
- ex:outputFormat
- ex:domainExpertise
- ex:historicalAccuracy

---

## 4. Expanded SHACL Shapes

Based on instructor feedback, new shapes are included.

### 4.1 Agent Shape (Revised)

- Must have name
- Must have  $\geq 1$  capability
- $\text{capabilityLevel} \in [0.0, 1.0]$
- Each `SpecializedAgent` must have  $\geq 1$  fallback (SPARQL constraint)

### 4.2 Task Shape (New)

- Must have required capability
- Must have complexity level in  $[0.0, 1.0]$
- Must map to at least one `TaskType`

### 4.3 Capability Shape (New)

- Must define capability type
- Must be referenced by at least one agent

### 4.4 Query Shape (New)

- Must have entity-extracted attributes
- Must route to exactly one agent

### 4.5 SPARQL Business Rules

Example:

```
# Every SpecializedAgent must have at least one fallback
ASK WHERE {
  ?agent a ex:SpecializedAgent .
  FILTER NOT EXISTS { ?agent ex:fallbackAgent ?fb . }
}
```

---

## 5. Required Cypher Queries (Updated for Neo4j)

Instructor requested operational queries; these are now rewritten in **Cypher**, optimized for Neo4j.

### Query 1: Select Best Agent by Capability Level

```
MATCH (agent:Agent)-[:HAS_CAPABILITY]->(cap:Capability),
      (task:Task)-[:REQUIRES_CAPABILITY]->(cap)
WHERE agent.capabilityLevel > 0.7
RETURN agent, agent.capabilityLevel AS capLevel
ORDER BY capLevel DESC;
```

### Query 2: Get Similar Agents for Fallback

```
MATCH (agent:Agent)-[:SIMILAR_TO]->(fallback:Agent)
RETURN fallback;
```

### Query 3: Retrieve Historical Routing Performance

```
MATCH (:RoutingDecision)-[:ROUTED_TO]->(agent:Agent)
WITH agent, avg(agent.confidence) AS avgConfidence
RETURN agent, avgConfidence
ORDER BY avgConfidence DESC;
```

### Query 4: Get Agents Matching TaskType Requirements

```
MATCH (taskType:TaskType)-[:REQUIRES_CAPABILITY]->(cap:Capability),
      (agent:Agent)-[:HAS_CAPABILITY]->(cap)
RETURN DISTINCT agent;
```

### Query 5: Retrieve Fallback Agent When Confidence Is Low

```
MATCH (agent:Agent)-[:FALLBACK_AGENT]->(fallback:Agent)
RETURN fallback;
```

## 6. Multi-Agent Workflow Architecture (Revised)

Multi-Agent Workflow Architecture (Revised) Instructor recommended clearer CrewAI-style decomposition.

### 6.1 QueryAnalyzer Agent

- Parses user query
- Extracts entities using GLiNER or GPT-4o-mini
- Produces structured representation:

```
{  
    "task_type": "code review",  
    "complexity": 0.6,  
    "domain": "technical",  
    "output_format": "summary"  
}
```

## 6.2 KnowledgeGraphQuery Agent

- Builds SPARQL/Cypher queries dynamically
- Retrieves candidate agents
- Ranks using:
  - capabilityLevel
  - historical performance
  - domain alignment

## 6.3 RoutingDecision Agent

- Chooses optimal agent
- Generates routing rationale
- Computes confidence score
- Handles tie-breaking logic

## 6.4 FeedbackCollector Agent

- Monitors success/failure
- Updates capabilityLevel, historicalAccuracy
- Allows the graph to improve over time

---

# 7. Entity Extraction Strategy (Revised)

Instructor asked for a firm commitment.

## Final Choice: GLiNER

Reasons: - Zero-shot extraction - High accuracy for custom entities - Fast inference - No need for large labeled datasets

Entities extracted: - TASK\_TYPE - COMPLEXITY\_LEVEL - DOMAIN - OUTPUT\_FORMAT

Mapping Example:

```
"code debugging" → ex:CodeDebuggingTask  
"complex" → complexityLevel 0.8
```

## 8. Agent Catalog (Demo Scope)

Instructor recommended 4–5 prototype agents.

- 1. WebSearchAgent**
- 2. CodeAnalysisAgent**
- 3. SummarizationAgent**
- 4. DataVisualizationAgent**
- 5. PerplexityFallbackAgent**

Each has: - capabilities - capabilityLevel - fallback relationships - sample historical performance

---

## 9. Demo Scenarios (Instructor Recommendation)

### **Scenario 1: “Find the latest research on LLM pruning.”**

Router → WebSearchAgent

### **Scenario 2: “Debug this Python snippet.”**

Router → CodeAnalysisAgent

### **Scenario 3: “Summarize this contract.”**

Router → SummarizationAgent

Extra:

### **Scenario 4: Low confidence**

Router → PerplexityFallbackAgent

---

## 10. Visualization & Explainability

Neo4j Bloom or GraphXR will show: - how agents relate - fallback chains - why the router selected an agent

Example explanation:

“Selected CodeAnalysisAgent because it has CodeUnderstanding capability (0.85), matches TaskType CodeDebuggingTask, and has the highest historical confidence (0.82).”

---

## 11. Final Implementation Plan

### Week 1

- Build RDF schema
- Implement full SHACL
- Create Neo4j graph with sample nodes

### Week 2

- Implement GLiNER extraction
- Implement QueryAnalyzer, KGQuery agents

### Week 3

- Implement router logic + confidence scoring
- Add fallback handling

### Week 4

- Build demo UI
  - Add feedback loop
  - Prepare elevator pitch & visualizations
- 

## 12. Elevator Pitch (Updated)

"The Smart Agentic Router solves the *agent selection problem* by using knowledge graphs to route user queries to the best specialized AI agent. Unlike simple keyword-based systems, our router uses semantic reasoning, agent profiling, and performance-based learning to deliver accurate, explainable, and continuously improving routing decisions."

---

## End of Updated Proposal